

Per-ROI hierarchical Bayesian GAM (projection away from HP vs distance from HP)

WHAT IS FIT (one model PER ROI – no joint master model):

```
proj ~ hsgp(distance, m=8, c=1.5) + (1 | subject)
```

- ``hsgp(...)`` is bambi's Hilbert-space Gaussian-process approximation, the equivalent of ``s(distance, k=8)`` in mgcv.
- ``(1 | subject)`` = per-subject random intercept.
- `family = "gaussian"`.
- HSGP length-scale prior:
 `ell ~ LogNormal(mu=0.693 = log(2.00°),
 sigma=0.50)`
 (slightly informative, centered around ~2.0°).
 HSGP marginal SD: `sigma ~ Exponential(1.0)`.
- Sampler: NUTS via PyMC, `draws = 2000`, `tune = 2000`,
 `chains = 4`, `target_accept = 0.99`.

DATA PIPELINE (same as `plot_projection_clusters.py`):

- `load_all_conditionwise(bids_folder='/data/ds-retsups', r2_thr=0.2)`.
- `filter_prf_inside_aperture(aperture_radius=3.17,
 max_fraction_outside=0.5,
 use_mean_only=False)`.
- Outlier clip: drop voxel-condition rows with
 `|proj_away_from_HP| > 1.5°` (`proj = -dy_rotated`).
- Per (subject, ROI, distance-bin): mean of `proj` across voxel-condition rows in the cell. Distance bin width = 0.5°, centers from 0.25° to 6.75°.

POSTERIOR SUMMARY:


- Smooth `f(distance)` is evaluated at 50 points from 0.50° to 6.50°. Per distance:
 - posterior median + 95% HDI of the population-level mean.
 - `P(f > 0)`: probability the smooth is > 0 at that distance.
 - `P(f < 0)`: probability the smooth is < 0 at that distance.
- A distance is flagged ("Bayesian cluster") when
 `P(>0) > 0.95` → red shading
 `P(<0) > 0.95` → blue shading

NOTES:

- Each ROI = a separate bambi model + separate posterior trace.
- Sampler warnings (divergences, etc.) are printed per-ROI.
- Default seed = 42; use `--seed` to vary.

ROIs fit: ['V1', 'V2', 'V3', 'V3AB', 'hV4', 'L0', 'T0', 'V0']


V1 — hierarchical Bayesian GAM on subject×bin means
n_subjects = 28, n_rows = 337, fit time = 185.2s, divergences = 59




$P(<0) > 0.95$ [0.50-1.72°]

$P(>0) > 0.95$ [2.46-3.44°]

V2 — hierarchical Bayesian GAM on subject×bin means
n_subjects = 28, n_rows = 335, fit time = 105.5s, divergences = 452




V3 — hierarchical Bayesian GAM on subject×bin means
n_subjects = 28, n_rows = 334, fit time = 98.8s, divergences = 338




proj. away from HP (deg)


V3AB — hierarchical Bayesian GAM on subject×bin means
n_subjects = 28, n_rows = 316, fit time = 100.7s, divergences = 61




hV4 — hierarchical Bayesian GAM on subject×bin means
n_subjects = 28, n_rows = 329, fit time = 78.6s, divergences = 1891




LO — hierarchical Bayesian GAM on subject×bin means
n_subjects = 28, n_rows = 325, fit time = 90.5s, divergences = 220



TO — hierarchical Bayesian GAM on subject×bin means
n_subjects = 28, n_rows = 290, fit time = 77.7s, divergences = 127



VO — hierarchical Bayesian GAM on subject×bin means
n_subjects = 28, n_rows = 276, fit time = 71.5s, divergences = 62



P(>0)>0.95 [2.21-3.07°]